

Cloud Storage

Dominic Duggan
Stevens Institute of Technology

1

Cloud Storage Models

- Relational Database
- Blob Storage
- NoSQL Storage
- Content Distribution Networks

2

Other Cloud Storage

- Queue Storage
 - See later
- File Storage
 - SMB file shares
 - OS API (internal)
 - REST API (cloud)

3

Cloud Storage Models

	Amazon Web Services	Windows Azure
Relational Database	Relational Database Service (RDS)	SQL Database
Blob Storage	Simple Storage Service (S3) Elastic Block Storage (EBS)	Blobs (Blocks) Blobs (Pages)
NoSQL Storage	Dynamo	Tables
Content Distribution	CloudFront	CDN

4



5

BLOB STORAGE

6

Blob Storage

- Blob: big list of binary data
 - Images, audio, video, etc
- Container: group of blobs
 - Flat file system
- Storage account
 - Max 100Tb
- Address:
`http://account.blob.core.windows.net/
container-name/blob-name`

7

Container Names

- Some rules for names
 - Valid DNS name
 - Letters, numbers, dash
 - *All letters must be lowercase*
- Root container
 - `http://image.blob.core.windows.net/upload/1.jpg`
 - `http://image.blob.core.windows.net/$root/1.jpg`
 - `http://image.blob.core.windows.net/1.jpg`

8

Blob Names

- Some rules for names
 - Case-sensitive
 - Resolved URL chars must be properly escaped
- Virtual directories:
 - landscape/grandcanyon.jpg
 - architecture/empirestate.jpg
 - architecture/eiffeltower.jpg
 - personality/jfk.jpg
 - http://image.blob.core.windows.net/upload/
personality/jfk.jpg

- Account name
- Container name
- Blob name


9

Types of Blobs

- Page
 - Optimized for random read/write
 - GET page (identified by offset)
 - Writes committed immediately
- Block
 - Optimized for streaming read
 - GET entire blob
 - Upload: 4Mb fragments & commit
- Append
 - Good for logging

10

Uploading Block Blob



© Redkar & Guidici

11

Endpoint Connection String

- Connection string in `web.config`:

```
<configuration>
  <appSettings>
    <add key="StorageConnectionString"
      value=
        "DefaultEndpointsProtocol=https;
        AccountName=AccountName;
        AccountKey=AccountKey" />
  </appSettings>
</configuration>
```

12

Accessing Blob Storage

- WCF Services API
- REST API
 - No client stubs
- StorageClient library
 - NuGet: install WindowsAzure.Storage

13

BLOB SERVICE REST API

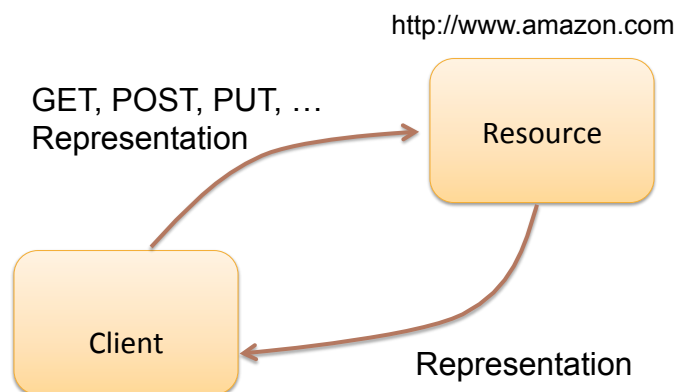
14

REST is not HTTP!

- ...but for this discussion we will assume it is
- Resources
 - Identified by URIs
- Representations
 - E.g. HTML files

15

Representational State Transfer



16

HTTP Request

GET /index.html HTTP/1.1

Host: www.example.org

...request headers...

17

HTTP Response

HTTP/1.1 200 OK

Date: Mon, 1 May 2011 21:38:14 GMT

Server: Apache/1.3.34 (Debian) mod_ssl/2.8.25
OpenSSL/0.9.8c ...

Last-Modified: Wed, 25 Nov 2009 12:27:01 GMT

ETag: "7496a6-a0c-4b0d2295"

Accept-Ranges: bytes

Content-Length: 2572

Content-Type: text/html

Via: 1.1 www.example.org

Vary: Accept-Encoding

...

18

Uniform Interface

- Retrieve: HTTP GET
- Create:
 - HTTP PUT for new URI or
 - HTTP POST for existing URI (server decides result URI)
- Modify: HTTP PUT, PATCH to existing URI
- Delete: HTTP DELETE

- Retrieve metadata only: HTTP HEAD
- Check which methods are supported: HTTP OPTIONS

- No other operations besides these

19

Account Operations

Operation	HTTP Verb	Cloud URI
List containers	GET	<code>http://acct.blob.core.windows.net?comp=list</code>

20

Properties vs Metadata

- System properties
 - Container: read-only
 - Blob: read-only and read-write
 - Ex: Etag, LastModifiedUTC, ContentEncoding, ContentType, ContentLanguage, etc
 - *blob.SetProperties()*
 - *blob.FetchAttributes()*
 - *blob.Properties.propName*
- User-defined metadata
 - *blob.SetMetadata()*
 - *blob.FetchAttributes()*
 - *blob.Metadata[key]*

21

Container Operations

Operation	HTTP Verb	Cloud URI
Create	PUT	<code>http://acct.blob.core.windows.net/container</code>
Get properties	GET/HEAD	<code>http://acct.blob.core.windows.net/container</code>
Set metadata	PUT	<code>http://acct.blob.core.windows.net/container?comp=metadata</code>
Get ACL	GET/HEAD	<code>http://acct.blob.core.windows.net/container?comp=acl</code>
Set ACL	PUT	<code>http://acct.blob.core.windows.net/container?comp=acl</code>
Delete container	DELETE	<code>http://acct.blob.core.windows.net/container</code>
List Blobs	GET	<code>http://acct.blob.core.windows.net/container?comp=list</code>

22

Blob Operations

Operation	HTTP Verb	Cloud URI
Write blob contents	PUT	<code>http://acct.blob.core.windows.net/container/blob</code>
Read blob contents	GET	<code>http://acct.blob.core.windows.net/container/blob</code>
Get properties	HEAD	<code>http://acct.blob.core.windows.net/container/blob</code>
Get metadata	GET/HEAD	<code>http://acct.blob.core.windows.net/container/blob?comp=metadata</code>
Set metadata	PUT	<code>http://acct.blob.core.windows.net/container/blob?comp=metadata</code>
Delete blob	DELETE	<code>http://acct.blob.core.windows.net/container/blob</code>

23

Blob Operations

Operation	HTTP Verb	Cloud URI
Put block	PUT	<code>http://acct.blob.core.windows.net/container/blob?comp=block&blockid=id</code>
Read block list (Block blob)	GET	<code>http://acct.blob.core.windows.net/container/blob?comp=blocklist&blocklisttype=[committed uncommitted all]</code>
Write block list (Block blob)	PUT	<code>http://acct.blob.core.windows.net/container/blob?comp=blocklist</code>
Copy blob	PUT	<code>http://acct.blob.core.windows.net/container/blob</code>
Lease blob	PUT	<code>http://acct.blob.core.windows.net/container/blob?comp=lease</code>

24

Blob Operations

Operation	HTTP Verb	Cloud URI
Snapshot	PUT	<code>http://acct.blob.core.windows.net/container/blob?comp=snapshot</code>
Write page (Page blob)	PUT	<code>http://acct.blob.core.windows.net/container/blob?comp=page</code>
Read page regions (Page blob)	GET	<code>http://acct.blob.core.windows.net/container/blob?comp=pagelist</code> <code>http://acct.blob.core.windows.net/container/blob?comp=pagelist&snapshot=datettime</code>

25

BLOB SERVICE: STORAGE CLIENT API

26

Accessing Blob Storage

- Accessing an account

```
using Microsoft.WindowsAzure;
using Microsoft.WindowsAzure.StorageClient;

CloudStorageAccount storageAccount =
    CloudStorageAccount
        .Parse(CloudConfigurationManager
            .GetSetting("StorageConnectionString"));

CloudBlobClient blobClient =
    storageAccount.CreateBlobClient();
```

27

Accessing Blob Storage

- Accessing a container

```
CloudBlobContainer container =
    blobClient
        .GetContainerReference("mycontainer");

CloudBlob blob =
    container
        .GetBlobReference("myblob");

@foreach (var blob in container.ListBlobs())
    <li> @blob.Uri </li>
```

28

Accessing Blob Storage

- Accessing a container in segments

```
CloudBlobContainer container =  
    blobClient  
        .GetContainerReference("mycontainer");  
  
ResultSegment<CloudBlob> resultSeg =  
    container.ListBlobsSegmented("",  
        BlobListingDetails.All, 25, null);  
  
while (resultSeg.ContinuationToken != null) {  
    foreach (var blob in resultSeg.Results) { ... }  
    resultSeg = resultSeg.GetNext(); }  
29
```

Accessing Blob Storage

- Accessing a blob: upload

```
CloudBlob blob =  
    container  
        .GetBlobReference("myblob");  
blob.Properties.ContentType = "image/jpeg";  
using  
    (var filestream =  
        System.IO.File.OpenRead(@"path\myfile"))  
    blob.UploadFromStream(fileStream);
```

30

Accessing Blob Storage

- Accessing a blob: download

```
CloudBlob blob =  
    container  
        .GetBlobReference("myblob");  
  
using  
    (var filestream =  
        System.IO.File.OpenWrite(@"path\myfile"))  
    blob.DownloadToStream(fileStream);
```

31

BLOB SERVICE ACCESS CONTROL

32

Types of Access

- Full public read access
 - Containers enumerable
- Public read access for blobs only
- Private
 - Only accessible to account owner
 - Authenticated with SHA256 HMAC
 - Authorization: [SharedKey|SharedKeyLite]
account-name:signature

33

Example Request

- List Prefixes in a Container

```
GET /upload?comp=list
    &delimiter=%2f
    &maxresults=100
    &timeout=30
    HTTP/1.1
x-ms-date: Sun, 30 Sep 2012 05:53:37 GMT
Authorization: SharedKey
    image:7euawYh5wNOGFJZGnvrn9vyR4y
Host: image.blob.core.windows.net
```

34

Example Response

```
HTTP/1.1 200 OK
Content-Type: application/xml
Server: Blob Service Version 1.0 Microsoft-HTTPAPI/2.0
x-ms-request-id: 7c490b17-8c99-43fa-ab8b-bde4cef032d7
Date: Sun, 30 Sep 2012 05:54:41 GMT
Content-Length: 408

<?xml version="1.0" encoding="utf-8"?>
<EnumerationResults
  ContainerName="http://image.blob.core.windows.net/upload">
  <MaxResults>100</MaxResults>
  <Delimiter>/</Delimiter>
  <Blobs>
    <BlobPrefix> <Name>Landscape/</Name> </BlobPrefix>
    <BlobPrefix> <Name>Architecture/</Name> </BlobPrefix>
```

35

Example Request

- List Blobs under a prefix in a Container

```
GET /upload?comp=list
    &prefix=Architecture%2f
    &delimiter=%2f
    &maxresults=100
    &timeout=30
HTTP/1.1
x-ms-date: Sun, 30 Sep 2012 05:57:24 GMT
Authorization: SharedKey
    image:E0V9XEPvs9J5zejM0HD+d3+3Lc2+B816HS9Vu2NwkaE
Host: image.blob.core.windows.net
```

36

Types of Access

- Setting public permissions

```
BlobContainerPermissions perm =  
    new BlobContainerPermissions();  
  
perm.PublicAccess =  
    BlobContainerPublicAccessType.Off;  
  
container.SetPermissions(perm);
```

37

Shared Access Signatures

- Delegate access via temporary URL (secret key)

```
PUT http://myacct.blob.core.windows.net/  
videos/myvideo.wmv?  
st=2012-10-21T05%3a52Z ← st: signed start (opt)  
&se=2012-10-31T08%3a49Z ← se: signed expiry  
&sr=c ← sr: signed resource (b or c)  
&sp=w ← sp: signed permission (r,w,d,l)  
&si=YWJjZGVmZw%3d%3d ← si: signed identifier (opt)  
&sig=Rcp6gPEaNGJAI$KAM%PIR$APANG%Ca%IL%O$V%E  
you%234so%m$uch2bqEArnfJxDgE%2bKH3TCChIs%3d  
HTTP/1.1 Host: myacct.blob.core.windows.net  
Content-Length: 19
```

38

Shared Access Signatures (1/2)

- Delegate access with temporary URL (delegator)

```
var sap =  
    new SharedAccessPolicy() {  
        Permissions = SharedAccessPermissions.Read  
            | SharedAccessPermissions.Write,  
        SharedAccessExpiryTime =  
            DateTime.UtcNow + TimeSpan.FromMinute(30)  
    };  
  
var sas =  
    container.GetSharedAccessSignature(sap);
```

39

Shared Access Signatures (2/2)

- Gain access with temporary URL (delegatee)

```
var sasCreds =  
    new StorageCredentialsSharedAccessSignature(sas);  
  
var client =  
    new CloudBlobClient  
        (storageAccount.BlobEndpoint, sasCreds);  
  
var secureBlob =  
    client.GetBlobReference ("mycontainer/myblob");
```

40

Signed Identifier (1/2)

- Define a *named* access policy

```
var perm = container.GetPermissions();

perm.SharedAccessPolicies.Add("myPolicy",
    new SharedAccessPolicy() {
        { Permissions = SharedAccessPermissions.Write }
    });

container.SetPermissions(perm,
    new BlobRequestOptions() {
        AccessCondition =
            AccessCondition.IfMatch(container.Properties.Etag)
    });
```

41

Signed Identifier (2/2)

- Associate the policy with a SAS

```
container.SetPermissions(perm, ...);

var sap =
    new SharedAccessPolicy() {
        SharedAccessExpiryTime =
            DateTime.UtcNow +
            TimeSpan.FromHours(24);
    };

var sas =
    container.GetSharedAccessSignature (sap,
        "myPolicy");
```

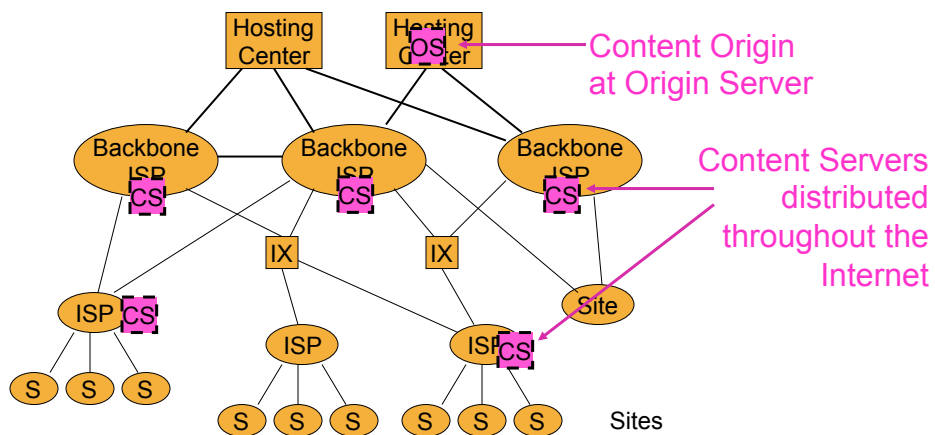
42

CONTENT DISTRIBUTION NETWORKS

43

Content Routing Principle


(a.k.a. Content Distribution Network)



44

Content Routing Principle


(a.k.a. Content Distribution Network)



Content is served from content servers nearer to the client

45


Cached CDN



1. Client requests content.

46


Cached CDN



1. Client requests content.
2. CS checks cache, if miss gets content from origin server.

47


Cached CDN



1. Client requests content.
2. CS checks cache, if miss gets content from origin server.
3. CS caches content, delivers to client.

48


Cached CDN



1. Client requests content.
2. CS checks cache, if miss gets content from origin server.
3. CS caches content, delivers to client.
4. Delivers content out of cache on subsequent requests.

49

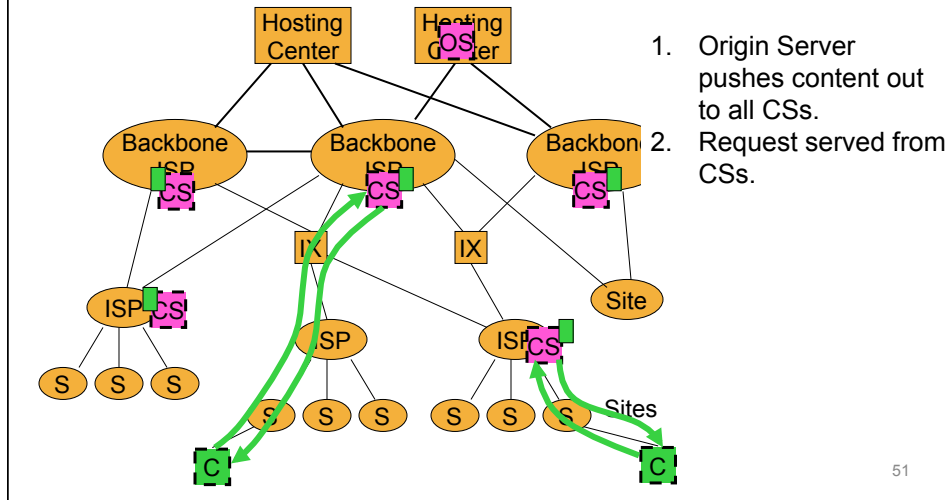
Pushed CDN



1. Origin Server pushes content out to all CSs.

50

Pushed CDN



CDN benefits

- Content served closer to client
 - Less latency, better performance
- Load spread over multiple distributed CSs
 - More robust (to ISP failure)
 - Flash loads spread over ISPs


52

CDN challenges

- Managing content distribution
 - Content lifetimes vs cache performance
 - Content synchronized and current
- Streaming content
 - Streaming vs file download?
 - Avoid proprietary protocols
- Live content
 - Television over IP

53

Streaming Content



54

Streaming Content

- Issue: Billing for content viewed
- Required:
 - Progressive download
 - e.g. Youtube
 - Throttling
- Solutions
 - Streaming HTTP
 - Adaptive Bit Rate (ABR) streaming

55

AZURE CDN

56

Azure CDN

- Blob storage account:
`http://account-name.blog.core.windows.net`
- URL for CDN access
`http://guid.vo.msecnd.net`
- Content must be public
- Content should be static
- Content may be streamed (audio, video)
- Cache endpoints based on demand

57

Caching Blobs in CDN

```
//Container must allow public access
BlobContainerPermissions perm =
    new BlobContainerPermissions();
perm.PublicAccess = BlobContainerAccessType.Container;
container.SetPermissions(perm);
```

```
//Set cache properties
blob.Properties.CacheControl = "public,max-age=30036000";
blob.SetProperties();
```

- Uncached URL:
 - `http://account.blob.core.windows.net/container/blob`
- Cached URL:
 - `http://guid.vo.msecnd.net/container/blob`

58

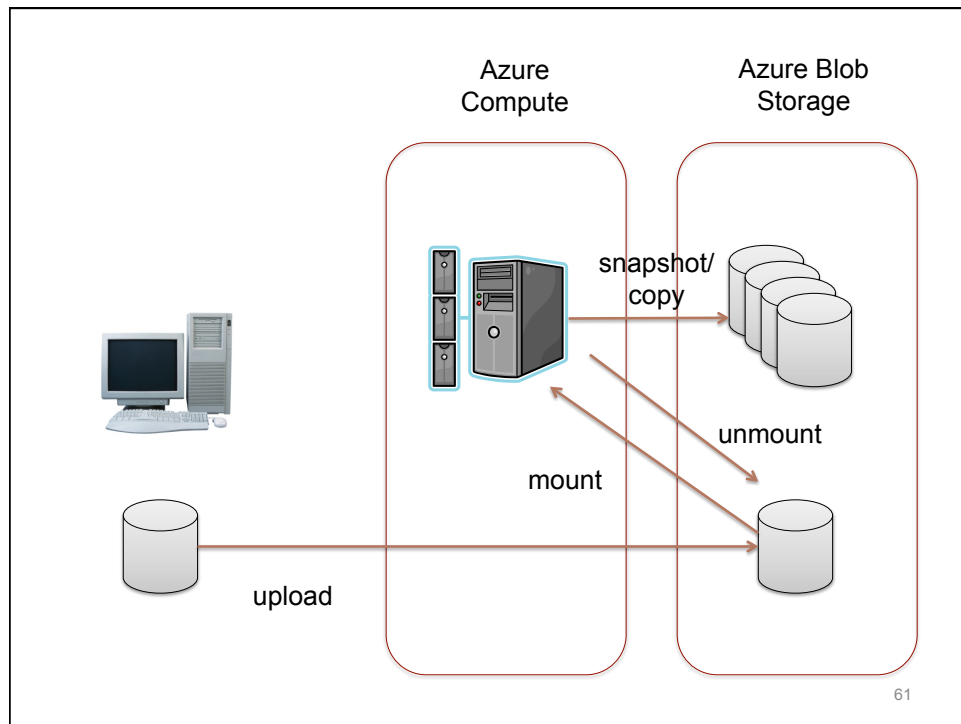
AZURE DRIVES

59

Azure Drives

- NTFS file system VHDs
- Stored as Page blobs
- Persist role instance state
 - `CloudDrive.Mount()`, `CloudDrive.Unmount()`
- Provide data for role instances
 - Multiple read-only mounts
- Create snapshots
 - Read-only: `CloudDrive.Snapshot()`
 - Make writable copies: `CloudDrive.Copy()`

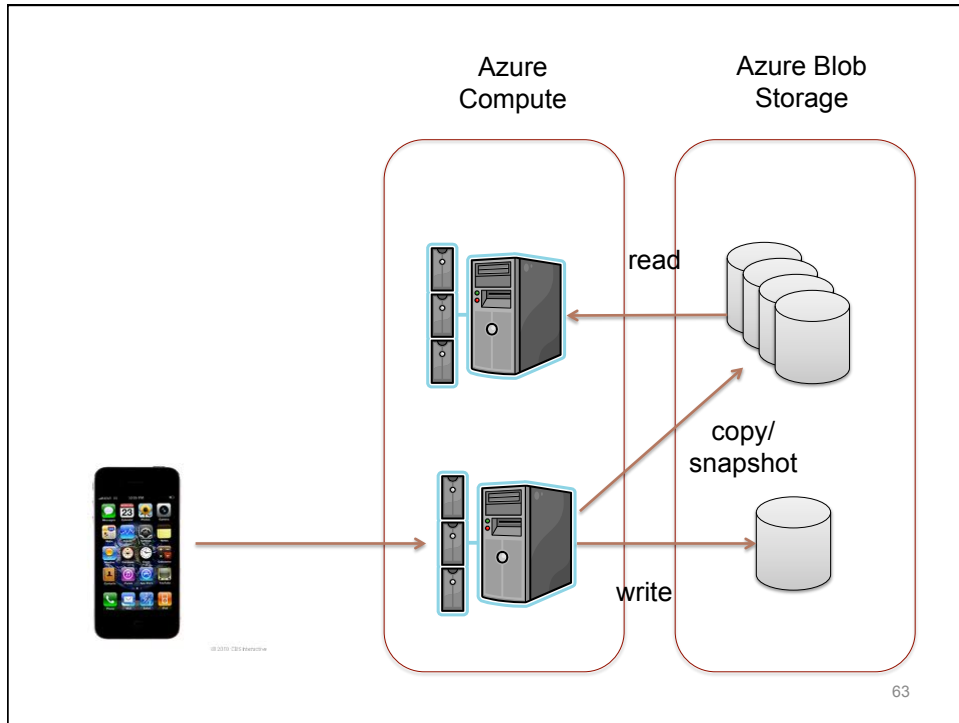
60



Azure Drive Scenarios

- Data storage for 3rd-party applications
- Read-only data storage for high-scale compute scenarios
 - E.g. Web app for data gathering
 - Data written to drives
 - Periodic snapshots
 - Worker role instances for processing

62



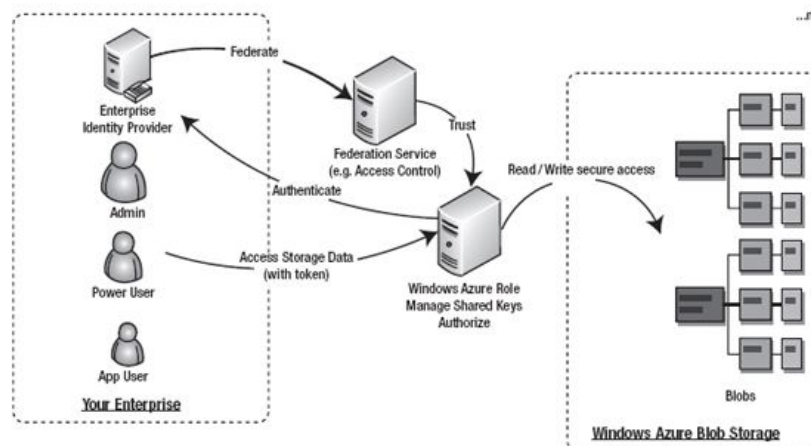
BLOB USAGE SCENARIOS

Blob Storage Scenarios

- Massive Data Uploads
 - Courier disks?
- Storage as a Service in the Cloud
 - Integrate with enterprise identity management
- Enterprise File Sync
 - Automatic backup

65


Storage as a Service



© Redkar & Guidici

66

Enterprise Sync



© Redkar & Guidici

67

Other Scenarios

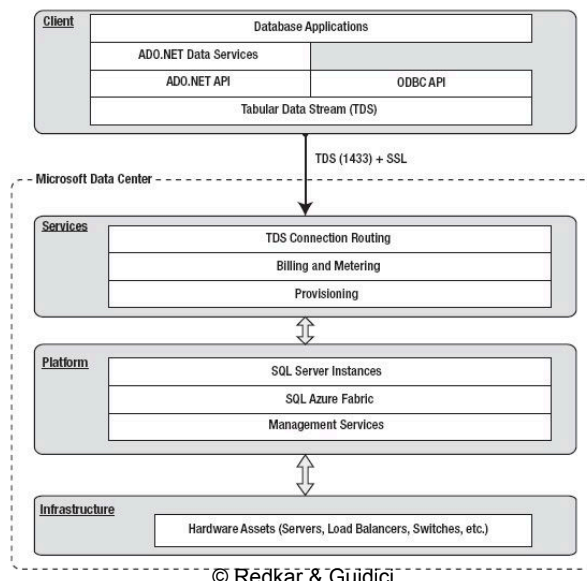
- Media Streaming
- Disaster Recovery Data Repository
- Application Store Repository
- Storing Mobile Apps and Data

68

SQL DATABASE

69

Architecture




70

© Redkar & Guidici

Platform Layer

- SQL server instances
 - Deployed databases
 - Replicas
 - OS instances
- Azure fabric
 - Replica set-up
 - Automatic failover
- Management service



71

Services Layer

- Gateway to platform layer
- Tabular data stream (TDS)
 - Port #1433 over SSL
 - DB provisioning (via Azure Fabric)
- Billing & metering
- Account provisioning

72

Client Layer

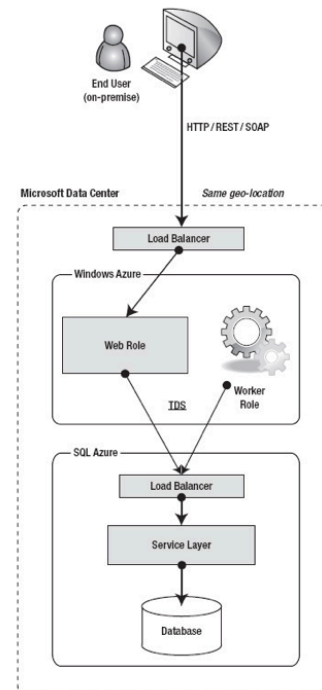
- Admin API
 - ADO.NET
 - ODBC
- TDS connection
 - Routed to primary

73

Access Pattern #1: Code Near

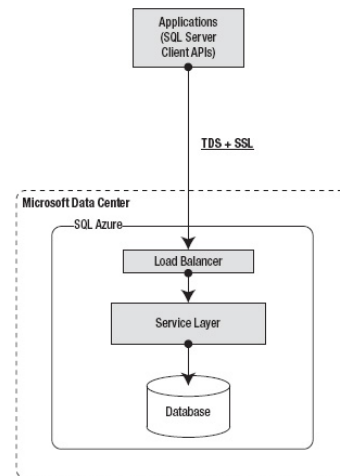
- Advantages:
 - Business logic close to DB
 - Latency
 - Open standards data access
 - HTTP, SOAP, etc
 - No lock-in
 - SQL Server client
- Disadvantage:
 - Performance
 - Azure as middle tier

© Redkar & Guidici



Access Pattern #2: Code Far

- Advantages:
 - Performance
 - Direct app connection
- Disadvantage:
 - Lock-in
 - TDS protocol
 - SQL Server clients
 - ADO.NET, ODBC, etc



© Redkar & Guidici

75

Connecting using ADO.NET

```
private string GetConnString() {  
    SqlConnectionStringBuilder sb =  
        new SqlConnectionStringBuilder();  
    string server =  
        "servername.ctp.database.windows.net";  
    sb.DataSource = server;  
    sb.InitialCatalog = "user database";  
    sb.Encrypt = true;  
    sb.TrustServerCertificate = true;  
    sb.UserID = "userName";  
    sb.Password = "password";  
    return sb.ToString();  
}
```

76

Connecting using ADO.NET

```
using (SqlConnection conn = new SqlConnection(GetConnectionString())) {  
    using (SqlCommand command = conn.CreateCommand()) {  
        conn.Open();  
  
        command.CommandText = "create table myTable" +  
                               "(Column1 primary key clustered, ...)";  
        command.ExecuteNonQuery();  
  
        command.CommandText = "select ... From MyTable";  
        using (SqlDataReader reader = command.ExecuteReader()) {  
            while (reader.read()) {  
                ...reader["Column1"].ToString().Trim()...  
            }  
        }  
    }  
}
```

77

DATABASE REVIEW

78

- SQL and Natural Join
- ```

select name, phone
from people
join phonenumber
on people.id=phonenumber.id

```

| Contacts |       |       |
|----------|-------|-------|
| id       | Name  | Phone |
| 1        | Joe   | 5532  |
| 1        | Joe   | 2234  |
| 1        | Joe   | 3211  |
| 2        | Jane  | 3421  |
| 3        | Chris | 2341  |
| 3        | Chris | 6655  |



| PhoneNumbers |    |       |
|--------------|----|-------|
| PhoneID      | Id | Phone |
| 1            | 1  | 5532  |
| 2            | 1  | 2234  |
| 3            | 1  | 3211  |
| 4            | 2  | 3421  |
| 5            | 3  | 2341  |
| 6            | 3  | 6655  |



| People |       |
|--------|-------|
| Id     | Name  |
| 1      | Joe   |
| 2      | Jane  |
| 3      | Chris |

| Name  | Phone |
|-------|-------|
| Joe   | 5532  |
| Joe   | 2234  |
| Joe   | 3211  |
| Jane  | 3421  |
| Chris | 2341  |
| Chris | 6655  |


79

## Transactions

Withdraw



Deposit



80



# Transactions



81

# Transactions



82

# Transactional Operations

amazon.com®

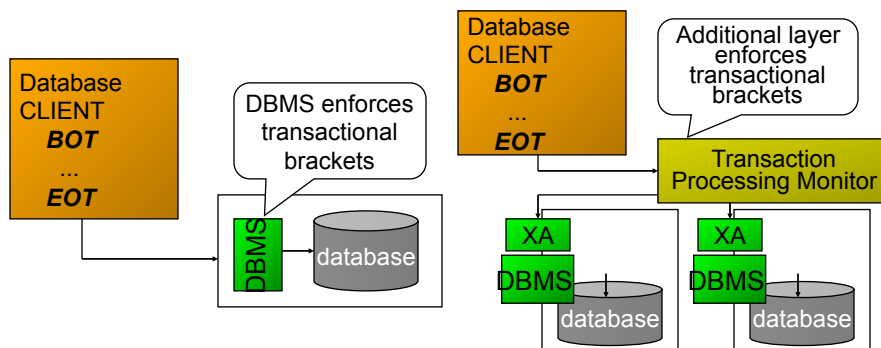
```
Begin Transaction
 AmtOwed = Amazon.Read (CustID);
 Balance = Paypal.Read (AccountID);
 if (AmtOwed ≤ Balance) {
 Paypal.Withdraw (AccountID, AmtOwed);
 Amazon.Deposit (CustID, AmtOwed);
 } else {
 Abort Transaction;
 }
End Transaction
```

PayPal®

83

# Distributed Databases

- Transactions and atomic commitment



84


## Relational Database Summary

- Database Schema
  - Normalized for efficiency
- SQL for ad-hoc queries
- Transactional updates
  - Atomtic
  - Consistent
  - Isolated
  - Durable

85

## Challenge: Big Data™

- Historical approach:  
*vertical scaling*
  - Limited
- Modern approach:  
*horizontal scaling*
  - *Sharding*
  - Azure: Federated SQL databases
  - Applications see data partitioning
  - No joins across partitions



86

# SQL vs NoSQL

## Relational

- Database Schema
  - Business data model
- SQL for ad-hoc queries
- ACID properties
  - Atomic
  - Consistent
  - Isolated
  - Durable

## NoSQL

- Unstructured
  - Web server logs
- Map-Reduce
- BASE properties
  - Basically Available
  - Soft state
  - Eventually consistent

87

## Fault-tolerance



88

## TABLE STORAGE

89

## Table Storage

- Property:
  - Name-value pair e.g. userid="JoeBlow"
- Entity: group of related properties
  - PartitionKey
    - Defines partition
  - RowKey
    - Defines order in a partition
  - Timestamp
  - Size limitations (1 Mb)
- Table:
  - Partitioned set of key-entity pairs

90

# Table Storage

- Property
- Entity
- Table

- Address:

`http://account.table.core.windows.net/Tables`

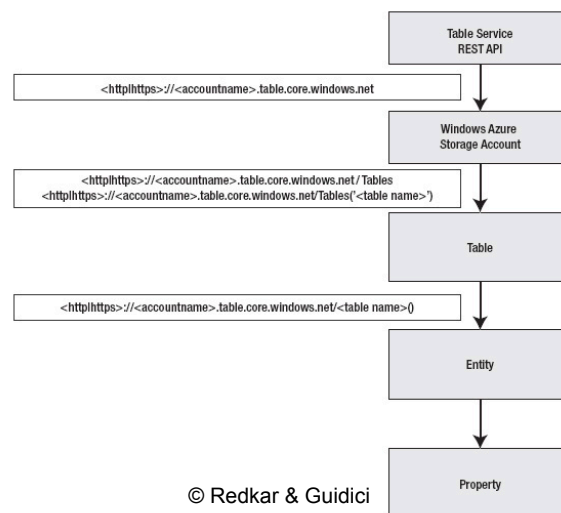
`http://account.table.core.windows.net/`

`Table('table-name')`

`http://acct.table.core.windows.net/table-name()`

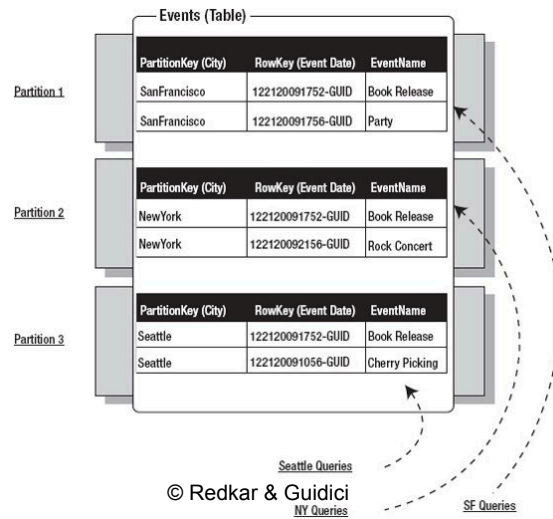
91

# Table Service Architecture



92

# Table Service Architecture



93

## Table Names

- Some rules for names
  - Valid DNS name
  - Letters, numbers
  - *All letters must be lowercase*
- Query parameters
 

```
http://dns-name/table-name()?$top=5
http://dns-name/table-name()?
 $filter=
 Date%20eq%20datetime'2012-09-30-20T00:00:00'
```

94

## Accessing Table Storage

- WCF Services API
- REST API
  - No client stubs
- StorageClient library
  - NuGet: install WindowsAzure.Storage

95

## TABLE SERVICE REST API

96



## Table Operations

| Operation    | HTTP Verb | Cloud URI                                                            |
|--------------|-----------|----------------------------------------------------------------------|
| Create table | POST      | <code>http://acct.table.core.windows.net/Tables</code>               |
| Query tables | GET       | <code>http://acct.table.core.windows.net/Tables()</code>             |
| Delete table | DELETE    | <code>http://acct.table.core.windows.net/Tables('table-name')</code> |

97

## REST Request

POST /Tables HTTP/1.1

```
User-Agent: Microsoft ADO.NET Data Services
x-ms-date: Sun, 9 Oct 2012 18:42:29 GMT
Authorization: SharedKeyLite
 image:pwFouPw+BPWzlaQPycII+K8zb+v6qygxZhp9fCdqRA=
Content-Type: application/atom+xml
Host: image.table.core.windows.net
<?xml version="1.0" encoding="utf-8" standalone="yes"?>
<entry xmlns:d=http://schemas.microsoft.com/ado/2007/08/dataservices
 xmlns:m=http://schemas.microsoft.com/ado/2007/08/dataservices/metadata
 xmlns="http://www.w3.org/2005/Atom">
 <title />
 <updated>2012-10-09T18:42:29.656Z</updated>
 <author> <name /> </author> <id />
 <content type="application/xml">
 <m:properties>
 <d:TableName>UserAccounts</d:TableName>
 </m:properties>
 </content>
</entry>
```

98

## REST Response

HTTP/1.1 201 Created

Content-Type: application/atom+xml; charset=utf-8

Location: http://image.table.core.windows.net/Tables('UserAccounts')

```
...
<?xml version="1.0" encoding="utf-8" standalone="yes"?>
<entry xmlns:d=http://schemas.microsoft.com/ado/2007/08/dataservices
 xmlns:m=http://schemas.microsoft.com/ado/2007/08/dataservices/metadata
 xmlns="http://www.w3.org/2005/Atom">
 <id>http://image.table.core.windows.net/Tables('UserAccounts')</id>
 <title />
 <updated>2012-10-09T18:42:29.656Z</updated>
 <author> <name /> </author>
 <link rel="edit" title="Tables" href="Tables('UserAccounts')"/>
 <content type="application/xml">
 <m:properties>
 <d:TableName>UserAccounts</d:TableName>
 </m:properties>
 </content>
</entry>
```

99

## Entity Operations

Operation	HTTP Verb	Cloud URI
Query entities	GET	<code>http://acct.table.core.windows.net/table-name()?\$filter=query-expr</code>
Insert entity	POST	<code>http://acct.table.core.windows.net/table-name</code>
Update entity	PUT	<code>http://acct.table.core.windows.net/table-name(PartitionKey="x", RowKey="y")</code>
Merge entity	MERGE	<code>http://acct.table.core.windows.net/table-name(PartitionKey="x", RowKey="y")</code>
Delete entity	DELETE	<code>http://acct.table.core.windows.net/table-name(PartitionKey="x", RowKey="y")</code>

100

## TABLE SERVICE: STORAGE CLIENT API

101

## Endpoint Connection String

- Connection string in web.config:

```
<configuration>
 <appSettings>
 <add key="StorageConnectionString"
 value=
 "DefaultEndpointsProtocol=https;
 AccountName=AccountName;
 AccountKey=AccountKey" />
 </appSettings>
</configuration>
```

102

## Accessing Table Storage

- Accessing an account

```
using Microsoft.WindowsAzure;
using Microsoft.WindowsAzure.StorageClient;

CloudStorageAccount storageAccount =
 CloudStorageAccount
 .Parse(CloudConfigurationManager
 .GetSetting("StorageConnectionString"));

CloudTableClient tableClient =
 storageAccount.CreateCloudTableClient();
```

103

## Accessing Table Storage

- Declaring an Entity

```
public class UserEntity : TableServiceEntity {

 public UserEntity
 (string lastName, string firstName) {
 this.PartitionKey = lastName;
 this.RowKey = firstName;
 }

 public UserEntity() { }

 public string Email { get; set; }
}
```

104

## Accessing Table Storage

- Adding Entities

```
TableServiceContext serviceContext =
 tableClient.GetDataServiceContext();

UserEntity user = new UserEntity("Joe", "Blow");
user.Email = "joe@blow.com";

serviceContext.AddObject("users", user);

serviceContext.SaveChangesWithRetries();
serviceContext.SaveChangesWithRetries
 (SaveChangesOptions.Batch);
```

105

## Accessing Table Storage

- LINQ query

```
CloudServiceQuery<UserEntity> query =
 (from e in serviceContext
 .CreateQuery<UserEntity>("users")

 select e)
 .AsServiceTableQuery<UserEntity>();

foreach (UserEntity e in query) {
 ...e.PartitionKey...e.Email...
}
```

106

## Accessing Table Storage

- LINQ query

```
CloudServiceQuery<UserEntity> query =
 (from e in serviceContext
 .CreateQuery<UserEntity>("users")

 select new { Email = e.Email })
 .AsTableServiceQuery<UserEntity>();

foreach (UserEntity e in query) {
 ...e.Email...
}
```

107

## Accessing Table Storage

- LINQ query

```
CloudServiceQuery<UserEntity> query =
 (from e in serviceContext
 .CreateQuery<UserEntity>("users")
 where e.PartitionKey == "Blow"
 && e.RowKey.CompareTo("M") < 0
 select e)
 .AsTableServiceQuery<UserEntity>();

foreach (UserEntity e in query) {
 ...e.PartitionKey...e.Email...
}
```

Ann Blow

Joe Blow

Nick Blow

108

## Accessing Table Storage

Ann Blow
<b>Joe Blow</b>
Nick Blow

- Single entity query

```
Query<UserEntity> singleQuery =
 (from e in serviceContext
 .CreateQuery<UserEntity>("users")
 where e.PartitionKey == "Blow"
 && e.RowKey == "Joe"
 select e);

UserEntity user1 = singleQuery.FirstOrDefault();
```

109

## Accessing Table Storage

- Update an entity

```
user1.Email = "joe.blow@yahoo.com";

serviceContext.UpdateObject(user1);

serviceContext.SaveChangesWithRetries();
```

110

## Accessing Table Storage

- Delete an entity

```
serviceContext.DeleteObject(user1);

serviceContext.SaveChangesWithRetries();
```

111

## Accessing Table Storage

- Update or insert an entity

```
UserEntity user2 = new UserEntity(...);

serviceContext.AttachTo("users", user2);

serviceContext.UpdateObject(user2);

serviceContext.SaveChangesWithRetries(
 SaveChangeOptions.ReplaceOnUpdate);
```

112



## Accessing Table Storage

- Delete without retrieving

```
UserEntity user2 = new UserEntity(...);

serviceContext.AttachTo("users", user2);

serviceContext.DeleteObject(user2);

serviceContext.SaveChangesWithRetries();
```

113

## Accessing Table Storage

- Deleting a table

```
CloudTableClient tableClient =
 storageAccount.CreateCloudTableClient();

tableClient.DeleteTableIfExists("users");
```

114

## **EXAMPLE: READER TRACKER**

115

### Reader Tracker

- Capture reader information and feedback
- Date when book purchased
- Where purchased
- New or used
- Personal info
- Feedback

116

## Reader Entity

- Define a class for reader entities:  
`public class Reader : TableServiceEntity { ... }`
- Properties:
  - Purchase date
  - Entry date
  - Country, State, City, Zip
  - Purchase location
  - Purchase type
  - Reader name
  - Reader URL
  - Feedback

117

## Reader Data Source

- `DataServiceContext`:
  - ADO.NET client-side state of invocations
- `ReaderDataSource`
  - Utility class for data queries
  - `CreateTable`, `DeleteTable`, `ListTables`
  - `AddReader`, `UpdateFeedback`
  - `Select`
  - `SelectByCity`, `SelectByCountry`,  
`SelectByState`
  - `SelectByPurchaseDate`
  - `SelectTopN`

118

## Keys for Reader Entity

- Find most dominant query in application
  - “Get all entities entered today.”
  - Query by date
  - Make PartitionKey to be EntryDate
  - All entries for same EntryDate in same partition
- Results must be ordered by time
  - RowKey based on EntryDate

119

## Keys for Reader Entity

```
public class DataServiceContext: TableServiceContext {
 public void AddRecord(
 DateTime purchaseDate,
 string country,
 string feedback, ...)
 {
 Reader reader = new Reader();
 reader.Country = country;
 reader.Feedback = feedback;
 reader.PurchaseDate = purchaseDate;
 ...
 this.AddObject("Reader", reader);
 this.SaveChanges();
 }
}
```

120

## Keys for Reader Entity

```
public class Reader : TableServiceEntity {
 public class Reader() { CreateKeys(); }

 public DateTime EntryDate { get; set; } ...

 public void CreateKeys() {
 EntryDate = DateTime.UtcNow;

 PartitionKey = EntryDate.ToString("mmddyyy");

 RowKey =
 string.Format("{0:10}_{1}",
 DateTime.MaxValue.Ticks - EntryDate.Ticks,
 Guid.NewGuid());
 }
}
```

121

## Reader Data Source

```
public class ReaderDataSource {

 private TableServiceContext context;
 public CloudStorageAccount Account { get; set; }
 public CloudTableClient TableClient { get; set; }
 public const string ENTITY_SET_NAME = "Reader";

 public ReaderDataSource
 (string storageAccountConnectionString) {
 Init(storageAccountConnectionString);
 context = TableClient.GetDataServiceContext();
 context.RetryPolicy =
 RetryPolicies.Retry(3,
 TimeSpan.FromSeconds(5));
 }
}
```

122

## Retry Policies

- Connection may be throttled
- Exception code 400,...,499, 501, 505: no retry
- Policies for other codes:
  - `RetryPolicies.NoRetry`
  - `RetryPolicies.Retry N`
  - `RetryPolicies.RetryExponential N` (default)

123

## Reader Data Source

```
public class ReaderDataSource {
 private void Init (string configName) {
 String connectionString =
 System.Web.Configuration
 .WebConfigurationManager.AppSettings(configName);
 Account =
 CloudStorageAccount.Parse(connectionString);
 ...
 }
}
```

```
<configuration>
 <appSettings>
 <add key="MyConnectionString"
 value="UseDevelopmentStorage=true" />
 </appSettings>
</configuration>
```

124

## Reader Data Source

```
public class ReaderDataSource {
 private void Init (string configName) {
 String connectionString =
 ConfigurationManager
 .ConnectionStrings[configName].ConnectionString;
 Account =
 CloudStorageAccount.Parse(connectionString);
 ...
 }
}
```

```
<configuration>
 <appSettings>
 <add key="MyConnectionString"
 value="UseDevelopmentStorage=true" />
 </appSettings>
</configuration>
```

125

## Reader Data Source


```
public class ReaderDataSource {
 private void Init (string configName) {
 String connectionString =
 RoleEnvironment
 .GetConfigurationSettingValue(configName);
 Account =
 CloudStorageAccount.Parse(connectionString);
 ...
 }
}
```

126

## Reader Data Source

```
public class ReaderDataSource {
 private void Init (string configName) {
 CloudStorageAccount
 .SetConfigurationSettingPublisher
 (MySettingPublisher1);
 Account =
 CloudStorageAccount.FromConfigurationSetting(configName);
 ...
 }

 private void MySettingPublisher1 (string configName,
 Func<string,bool> configSettingPublisher) {
 String connectionString =
 RoleEnvironment
 .GetConfigurationSettingValue(configName);
 configSettingPublisher(connectionString);
 }
}
```




127

## Reader Data Source

```
public class ReaderDataSource {
 private void Init (string configName) {
 CloudStorageAccount
 .SetConfigurationSettingPublisher
 (MySettingPublisher2);
 Account =
 CloudStorageAccount.FromConfigurationSetting(configName);
 ...
 }

 private void MySettingPublisher2 (string configName,
 Func<string,bool> configSettingPublisher) {
 String connectionString =
 ConfigurationManager
 .ConnectionStrings[configName].ConnectionString;
 configSettingPublisher(connectionString);
 }
}
```



128



## Reader Data Source

```
public class ReaderDataSource {
 private void Init (string configName) {
 if (RoleEnvironment.IsAvailable) {
 CloudStorageAccount
 .SetConfigurationSettingPublisher
 (MySettingPublisher1);
 } else {
 CloudStorageAccount
 .SetConfigurationSettingPublisher
 (MySettingPublisher2);
 }

 Account =
 CloudStorageAccount
 .FromConfigurationSetting(configName);
 }
}
```

129

## Reader Data Source

```
public class ReaderDataSource {
 private void Init (string configName) {
 ...
 Account = CloudStorageAccount
 .FromConfigurationSetting(configName);

 TableClient = Account.CreateCloudTableClient();
 TableClient.RetryPolicy =
 RetryPolicies.Retry(3,
 TimeSpan.FromMilliseconds(100);
 }
}
```

130

## Reader Data Source

```
public class ReaderDataSource {

 public IEnumerable<Reader> Select() {
 var results =
 from g
 in context.CreateQuery<Reader>(ENTITY_SET_NAME)
 where g.PartitionKey ==
 DateTime.UtcNow.ToString("mmddyyyy")
 select g;
 var r = results.ToArray<Reader>();
 return r;
 }
}
```

131

## Reader Data Source

```
public class ReaderDataSource {

 public IEnumerable<Reader> SelectByCity(string city) {
 var results =
 from g
 in context.CreateQuery<Reader>(ENTITY_SET_NAME)
 where g.PartitionKey ==
 DateTime.UtcNow.ToString("mmddyyyy")
 && g.City == city
 select g;
 var r = results.ToArray<Reader>();
 return r;
 }
}
```

```
GET http://myacct.table.core.windows.net/Reader?
$filter=City%20eq%20...
```

132

## Reader Data Source

```
public class ReaderDataSource {

 public IEnumerable<Reader> SelectTopN (int TopNum) {
 var results =
 (from g
 in context.CreateQuery<Reader>(ENTITY_SET_NAME)
 where g.PartitionKey ==
 DateTime.UtcNow.ToString("mmddyyyy")
 select g)
 .Take(topNum);

 var r = results.ToArray<Reader>();
 return r;
 }

 GET http://myacct.table.core.windows.net/Reader?
 $top=N
```

133

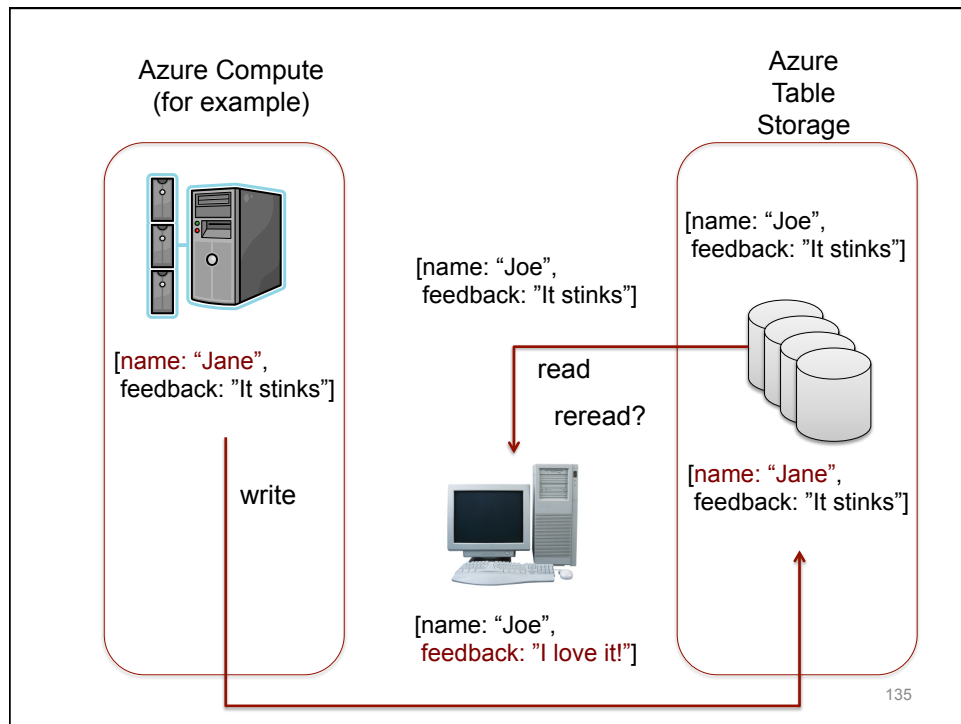
## Reader Data Source

```
public class ReaderDataSource {

 public void UpdateFeedback (string partitionKey,
 string rowKey,
 string feedback) {

 var results =
 from g
 in context.CreateQuery<Reader>(ENTITY_SET_NAME)
 where g.PartitionKey == partitionKey
 && g.RowKey == rowKey
 select g;
 var e = results.FirstOrDefault<Reader>();
 e.Feedback = feedback;
 context.MergeOption = MergeOption.PreserveChanges;
 context.SaveObject(e);
 context.SaveChanges();
 }
}
```

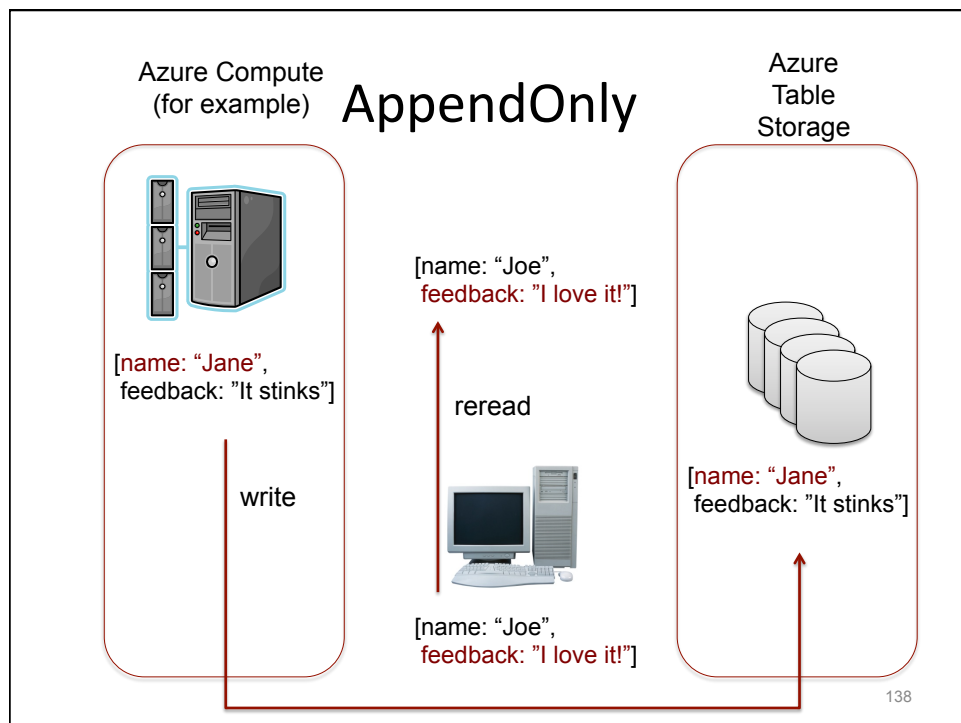
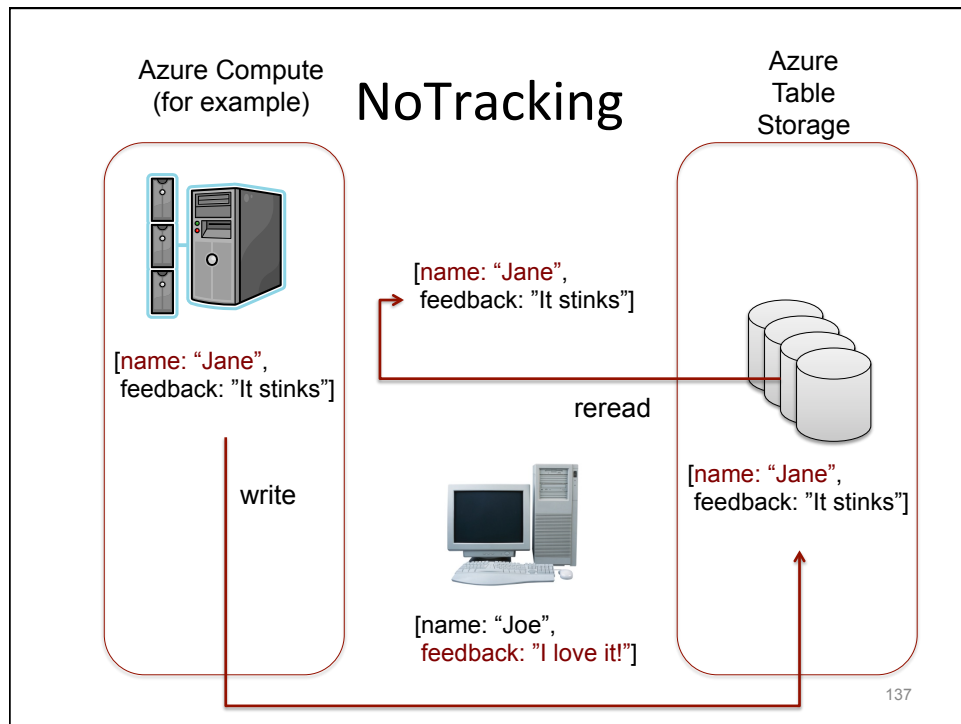
134

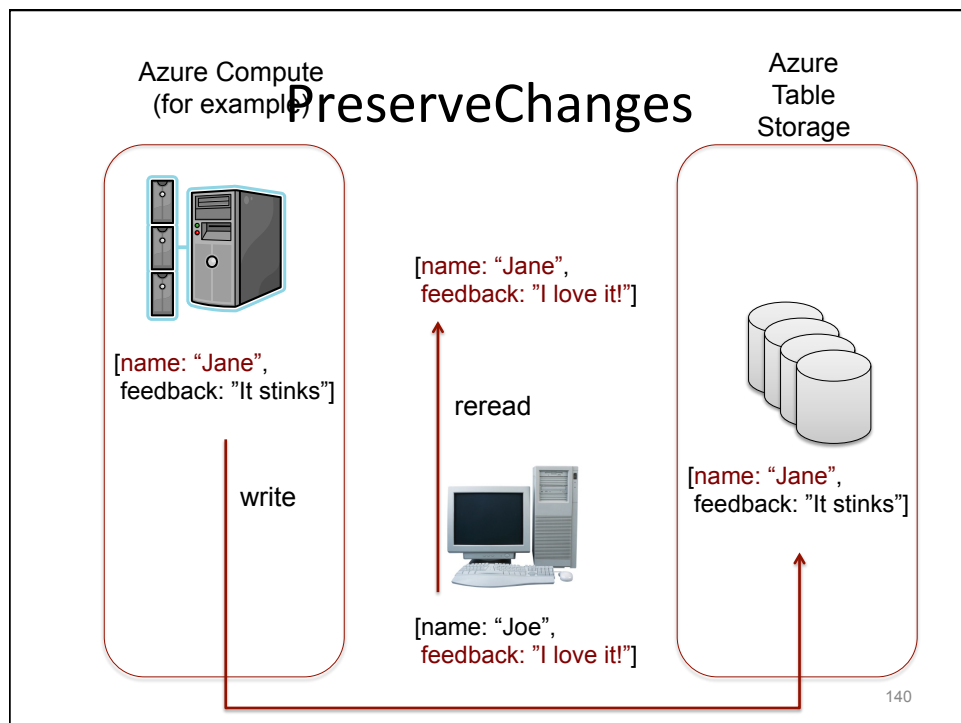
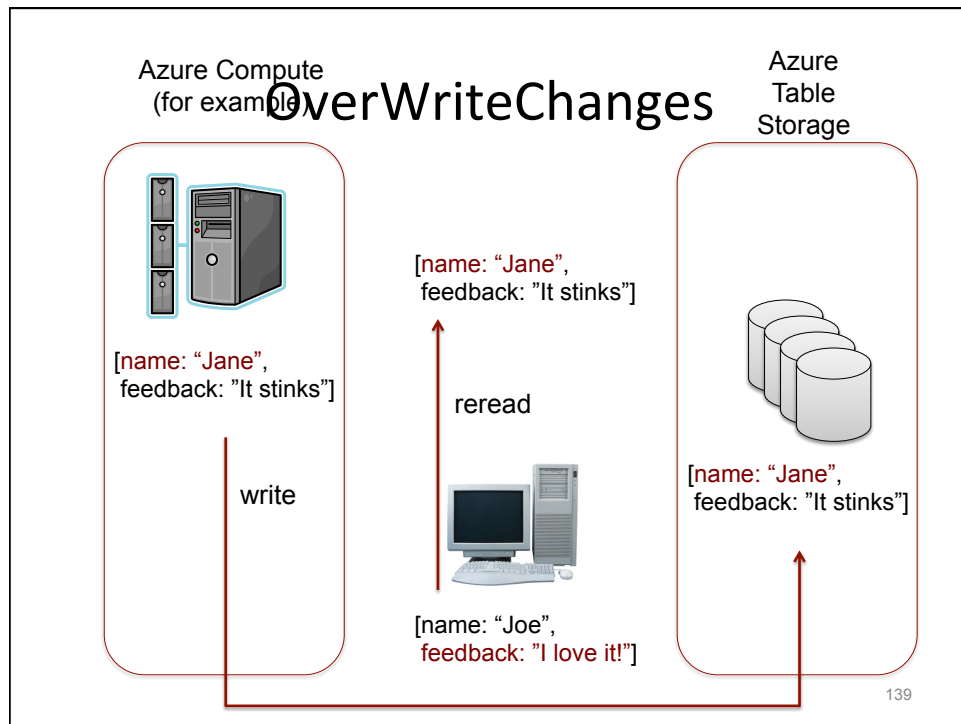


## Merge Options

- Merge: Update properties without replacing entity
- NoTracking
  - Always loaded from server
- AppendOnly (default)
  - Add to local cache
- OverwriteChanges
  - Server values take precedence
- PreserveChanges
  - Change etag with server etag, preserve local changes

136





## TABLES USE CASES

141

## Table vs Relational DB

- Migration
- Cost
  - Table much cheaper
  - Even when not normalized
- Lock-in
- Utility Tables
  - No relational requirements
  - Ex: e-commerce requirements
  - Ex: User profiles
  - Ex: Tracing information

142

## Table Use Case

- Storing Performance Counters
  - Ex: logs from Compute instances

143

## Storing Performance Counters

```
public class PerformanceData :
 TableServiceEntity
{
 public Int64 EventTickCount { get; set; }
 public string DeploymentId { get; set; }
 public string Role { get; set; }
 public string RoleInstance { get; set; }
 public string CounterName { get; set; }
 public double CounterValue { get; set; }
}
```

144



## Querying Performance Data

```
var account =
 new CloudStorageAccount(
 new StorageCredentialsAccountAndKey
 (storageAccountName, storageKey),
 true);

var context =
 new PerformanceDataContext
 (account.TableEndpoint.ToString(),
 account.Credentials);

var data = context.PerfData;
```

145

## Querying Performance Data

```
var data = context.PerfData;

DateTime tf =
 DateTime.UtcNow.Subtract(
 TimeSpan.FromMinutes(timeFrameInMinutes));

List<PerformanceData> selectedData =
 (from d in data
 where (DateTime.Compare(tf, d.Timestamp) < 0)
 select d)
 .ToList<PerformanceData>();
```

146

## Querying Performance Data

```
var data = context.PerfData;

DateTime tf =
 DateTime.UtcNow.Subtract(
 TimeSpan.FromMinutes(timeFrameInMinutes));

List<PerformanceData> selectedData =
 (from d in data
 where (DateTime.Compare(tf, d.Timestamp) < 0)
 select d)
 .AsTableServiceQuery<PerformanceData>();
```

147